

Problem 1

(Max flow using the bottleneck path) [REDACTED]

Problem 2

(Push-Relabel algorithm) The algorithms explored up to this point share a fundamental attribute: they aim to preserve a conservation of flow constraints while working towards severing the connection between s and t (aka a “cut”). But in the 1980s, some researchers began thinking about max-flow another way: what if, instead of maintaining a connection between s and t , we start with them already disconnected and then “relax” our flow to maintain conservation? This paradigm shift led to the discovery of the Push-Relabel algorithm, which we now present below.¹

Before we do so, we need to introduce a few definitions.

- A **preflow** is a nonnegative vector $\{f_e\}_{e \in E}$ that satisfies two constraints:
 1. **Capacity constraints:** $f_e \leq c_e$ for every edge $e \in E$;
 2. **Relaxed conservation constraints:** for every vertex v other than s ,
amount of flow entering $v \geq$ amount of flow exiting v .

We call the difference between the LHS and the RHS the **excess** of a vertex.

- The **height** of a vertex $h(v)$ is defined by a function $h: V \rightarrow \mathbb{N}$. Intuitively, you can think of h as a potential function that flows downhill from s to t .

Now we can introduce our algorithm, defined on a graph $G = (V, E, c_e, s, t)$. Also denote the residual graph as $G^f = (V, E^f, c_e^f, s, t)$.

¹For additional motivation and context behind this algorithm, please consult [this](#) lecture by Tim Roughgarden.

Algorithm 1 PUSHRELABELMAXFLOW(G)

Initialize $h(s) = n$

Initialize $h(v) = 0$ for all $v \neq s$

Initialize $f_e = c_e$ for all $e \in$ outgoing edges from s

Initialize $f_e = 0$ for all $e \notin$ outgoing edges from s

while exists vertices $\mathcal{V} \in V \setminus \{s, t\}$ with excess flow > 0

 Choose vertex v such that $v = \arg \max_{\nu \in \mathcal{V}} h(\nu)$.

if exists outgoing edge (v, w) in E^f such that $h(v) = h(w) + 1$ **then**

 PUSH(v, w)

else

 RELABEL(v)

▷ Subroutine for Push operation

procedure PUSH(v, w)

$\Delta = \min \left\{ \text{excess flow}(v), c_{(v,w)}^f \right\}$

 Push Δ units of flow along (v, w)

▷ Subroutine for Relabel operation

procedure RELABEL(v)

 Increment $h(v)$ by 1


- (a) The network below represents the (hypothetical) connections between a number of international airports, where the capacity of an edge represents the maximum number of flights that can be scheduled between two airports on a given day.

Imagine an urgent situation, like a massive concert or a sports event suddenly announced in Atlanta, has just caused a surge of travelers needing to fly from Los Angeles (LAX) to Atlanta (ATL) ASAP. The big question is: how many flights can actually make this journey in a single day without overstretching the system?

Run the Push-Relabel Algorithm to answer this question. Along with your final numerical answer, include the following:

- i. A sketch of the final residual graph with the accompanying capacities (indicate edges in $E^f \setminus E$ with dashed lines), as well as the final heights of each node.
- ii. The number of times the **Push** subroutine has been called, as well as the number of times the **Relabel** subroutine was called.

In the case of ties, choose the node with the smallest lexicographical order.



- (b) A basic implementation of the algorithm in the exposition gives us a runtime of $\mathcal{O}(mn^2)$. Give an implementation of this algorithm that runs in $\mathcal{O}(n^3)$ time, and prove the runtime of your algorithm.

Hints: first prove the running time bound assuming that, in each iteration, you can identify the highest vertex with positive excess in $\mathcal{O}(1)$ time. The hard part is to maintain the vertices with positive excess in a data structure such that, summed over all of the iterations of the algorithm, only $\mathcal{O}(n^3)$ total time is used to identify these vertices.

- (c) (Optional) It can be shown (with a little more work) that a tight bound for this algorithm is $\Theta(n^2\sqrt{m})$ ². Although this bound is nothing special in the worst-case (for example, Dinic's algorithm runs in time $\mathcal{O}(nm^2)$), the PR algorithm is often the one used.³ What are some reasons why this algorithm may perform better in practice?

²Asymptotically, a \sqrt{m} term is better than a n term.

³This dichotomy with theoretical bounds and practical application is oftentimes the case with flow-based networks; for example, in [computer vision](#).